

DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Empirical Evaluation of Test Suite
Reduction**

Adrian Regenfuß

DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Empirical Evaluation of Test Suite
Reduction**

**Empirische Evaluation von Test-Suiten
Reduktion**

Author:	Adrian Regenfuß
Supervisor:	Prof. Dr. Dr. h.c. Manfred Broy
Advisor:	Dr. Elmar Jürgens, Raphael Nömmner, Roman Haas
Submission Date:	Submission date

I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

Munich, Submission date

Adrian Regenfuß

Acknowledgments

Abstract

As a response to ever-growing test suites with long runtimes, several different approaches have been developed to reduce the time for test suite execution to give useful results. One of these approaches is test suite reduction: selecting a sample of tests that maximizes coverage on the tested source code. This paper attempts to replicate the findings in [Cru+19], which borrows techniques from big data to handle very large test suites. We use independently generated testing data from open source projects, and find that TODO.

Contents

Acknowledgments	iii
Abstract	iv
1 Introduction	1
2 Terms and Definitions	3
3 Related Work	4
3.1 Test Case Priorization	4
3.2 Test Case Selection	5
3.3 Test Suite Reduction	5
3.3.1 Greedy Selection	6
3.3.2 Clustering	6
3.3.3 Searching	7
3.3.4 Hybrids and Other Methods	7
4 Approach	8
4.1 Replicating "Scalable Approaches for Test Suite Reduction"	8
4.2 Implemented Algorithms	8
4.2.1 FAST	8
4.2.2 Random Selection	12
4.2.3 Adaptive Random Testing	12
4.2.4 Greedy Algorithm	14
5 Replication Study	16
5.1 Research Questions	16
5.2 Study Design	17
5.3 Study Objects	17
5.4 Selecting Projects	17
5.5 Study Setup	18
5.5.1 Combining Tests Suites	18
5.5.2 Generating Coverage Information	19

Contents

5.5.3	Collecting Fault Coverage Information	19
5.6	Results	19
5.6.1	Research Questions	20
5.7	Discussion	22
5.8	Threats to Validity	22
5.8.1	Construct Validity	23
5.8.2	Internal Validity	23
5.8.3	Conclusion Validity	24
5.8.4	External Validity	24
6	Future Work	25
7	Summary	26
	List of Figures	27
	List of Tables	28
	Bibliography	29

1 Introduction

3 pages

Software often has faults: ways in which the actual behavior of the software diverges from the intended or specific behavior. Computer scientists have devised different strategies for finding and removing bugs: Formal software verification, code reviews, and different types of testing: unit testing, which tests the behavior of small software modules (such as classes), model-based testing, which automatically generates test suites from the specification of a software system, and regression testing, which shows developers whether they have introduced new bugs after the last change to the software.

Regression testing takes up a significant portion of development cost, and the resulting test suites can grow quite significantly in size and execution time, which hinders development speed and increases costs.

Similarly, model-based testing suffers from issues of combinatorial explosion, and the resulting test suites tend to be very big.

To mitigate the costs and runtimes of regression and model-based test suites, different strategies have been devised: test case selection, which selects a subset of tests for the current execution of tests, test suite reduction, which permanently deletes a subset of tests from the test suite, and test case prioritization, which changes the order of test execution to maximize the amount of faults that is found early in the test suite execution.

[Cru+19] presents a new family of test suite reduction algorithms, called the FAST algorithms (first developed in [Mir+18]), and compares them to the algorithms presented in [Che+10], as well as the greedy algorithm presented in [Rot+01]. They implement 4 algorithms from the FAST family, 2 from the ART family and the greedy algorithm, and compare the different algorithms on 10 different test programs and their test suites. They compare the performance of the test suite reduction algorithms on 3 different variables: test suite reduction, fault detection loss, and runtime.

This work attempts to replicate their findings using the data from 6 additional open-source projects, as well as adding random test case selection as a baseline test suite reduction method to compare the other methods to (following **Recommendation 8** from [Kha+18]).

This work then attempts to determine how different test suite reduction strategies

compare to each other in terms of time performance, fault detection loss and magnitude of reduction.

2 Terms and Definitions

This paper uses some terms and abbreviations that are common in the literature on test suite reduction, but not used much elsewhere.

The System Under Test (SUT) is a software system being tested using a test suite.

The Test Suite Reduction (TSR) is a metric that evaluates the magnitude of the reduction of a test suite. Given a test suite $T = \{t_1, \dots, t_n\}$ and a reduced test suite $T' \subset T$, the TSR of the test suite is determined using the formula

$$\text{TSR}(T, T') = \frac{|T| - |T'|}{|T|}$$

Everything else equal, higher values for TSR are better.

The abbreviation TSR also sometimes stands for the technique of test suite reduction, i.e. the permanent removal of redundant test cases from a test suite. In this text, "TSR" will only stand for the metric, and the technique will be written out as "test suite reduction".

The Fault Detection Loss (FDL) is a metric that evaluates the amount of faults detected by a reduced test suite. Let F be the faults detected by T , and F' be the faults detected by T' , then the fault detection loss is calculated as such:

$$\text{FDL}(F, F') = \frac{|F| - |F'|}{|F|}$$

In general, lower FDL values are better. A related metric (e.g. used in [Kha+16] or [YH12]) is the Fault Detection Effectiveness (FDE), defined as $1 - \text{FDL}(F, F')$.

3 Related Work

5 pages

The literature on handling large test suites is comprehensive, and this summarization will only cover small parts of it. For an older, but comprehensive overview dividing the field into test-suite reduction, test case selection and test case prioritization, see [YH12]. A newer overview focused exclusively on test suite reduction is [Kha+16], which attempts to create a full taxonomy of TSR frameworks/tools and specifically their implementations. [Kha+18] focuses on the algorithmic approach (i.e. greedy/clustering/searching/hybrid) of the TSR system and gives recommendations for future the evaluation of TSR methods.

The methods for improving the runtimes of large test suites discussed in the literature can be divided into three related approaches ([YH12]): Test suite minimization (here called test suite reduction), test case prioritization, and test case selection.

Let $T = \{t_1, t_2, \dots, t_n\}$ be a set of n tests, and a set of requirements $R = \{r_1, \dots, r_m\}$. These requirements can take very different forms: they can correspond to coverage of lines/branches/functions, parts of an explicit specification that must be tested, runtimes of individual tests, or faults that were discovered in the past.

Let $m : T \rightarrow \mathcal{P}(R)$ a function that maps test cases to the requirements they test (one test can test multiple requirements, but any requirement needs only one test case to be fulfilled, i.e. there can be no requirements that need 2 or more tests to be fulfilled).

3.1 Test Case Prioritization

Test case prioritization attempts to find a permutation $T' \in S_T$ (S_T being the set of all permutations of T) of test cases that cover as many requirements as early as possible.

More formally, the test case prioritization problem is to find a permutation $T' \in S_T$ so that

$$\forall T'' \in S_T : T' \neq T'' \wedge \forall i \in 1..n : \left| \bigcup_{j=1}^i m(T'(j)) \right| \geq \left| \bigcup_{j=1}^i m(T''(j)) \right|$$

Test case prioritization subsumes test suite reduction: Given an ordering T' of test cases, it is easy to select the first B test cases in the adequate scenario or the first i test

cases that together fulfill all requirements.

3.2 Test Case Selection

Test case selection attempts to temporarily find a subset of tests that maximize the requirements for a set of recent changes to the software.

[YH12] formulate the problem of test case selection as

"*Given:* The program, P , the modified version of P , P' and a test suite, T .

Problem: Find a subset of T , T' , with which to test P' ."

Test case selection usually assumes a high degree of transparency and the availability of a lot of information about the SUT, such as execution graphs, coverage information, and sometimes information about the running times of individual tests.

3.3 Test Suite Reduction

Test suite reduction, on the other hand, attempts to permanently reduce the size of the test suite by removing redundant tests. The distinguishing feature from test case selection is that it doesn't take into account recent changes.

Generally, in test case reduction two different cases are distinguished: adequate and inadequate reduction.

In adequate test suite reduction, the goal is to find the smallest subset of T so that all R are satisfied:

$$T_r = \operatorname{argmin}_{T_r} |T_r \subset T : R = \bigcup_{t_r \in T_r} m(t_r)|$$

This goal, however, is rarely fulfilled. Often, the objective is to keep the resulting test suite as small as possible, since reaching this optimum is computationally hard (as researchers have remarked ([Kha+16]), it is equivalent to the NP-complete set cover problem ([GJ79])).

Instead, researchers attempt to reduce the size of test suites further, while still keeping running times small.

Other possible goals include to reduce the total runtime of the reduced test suite, but due to the absence of information about test runtimes, this goal has been studied less.

Adequate reduction needs the requirements R to be present at the time of reduction to judge whether all requirements have been fulfilled.

The second case of test suite reduction is the inadequate one. An inadequate test suite reduction fixes the size of the reduce test suite to a budget $B = |T_r|$. This cannot

guarantee that all requirements are completely fulfilled, but often reduces runtimes significantly.

Inadequate reduction can work without access to the requirements. For example, the FAST family needs only the content of the testcases to perform a reduction. However, inadequate reduction gives no guarantees about the performance of the resulting reduced test suite.

[Kha+18] distinguishes 4 different types of test suite reduction approaches: Greedy selection, Clustering, Searching, and Hybrid methods.

3.3.1 Greedy Selection

Greedy approaches, first introduced by [Rot+01] (in the context of test case prioritization), work by selecting test cases based on test requirements $R = \{r_1, \dots, r_m\}$ (mostly statement/branch/function coverage).

[Kha+18] describe the fully general case of greedy-based test suite reduction, in which a test case is first selected using the coverage criterion and a heuristic, and then removed from the test set and added to the set of the reduced test suite, until the coverage criterion is completely fulfilled by the reduced test suite.

Two simple heuristics, described more in detail later, are the total heuristic (choosing the test with the highest amount of coverage from the test set) and the additional heuristic (choosing the test with the highest amount of coverage yet not covered by the reduced test suite).

Greedy approaches to test suite reduction can also include approaches based on integer linear programming, such as in

3.3.2 Clustering

Clustering approaches to test suite reduction attempt to maximize a distance metric on the resulting test suite. This is achieved by separating T into disjunctive subsets using a clustering algorithm so that the difference between two subsets doesn't fall under a given threshold, and then sampling test cases from the subsets (in the optimal case one representative test case per subset).

This method doesn't need the test requirements to be available: many approaches (for example from the FAST family examined in this paper [Mir+18]) calculate the similarity of two test cases using the string of the test case's source code. However, those approaches can only be inadequate: The test case strings give no information about when the requirements are completely fulfilled.

Examples of clustering-based approaches to test suite reduction are of course the FAST algorithms from [Cru+19] and [Mir+18], but also the ART family from [Che+10]

and CANDO.

3.3.3 Searching

Search-based methods for test suite reduction evaluate subsets $S = \{s_1, \dots, s_m\} \subset \mathcal{P}(T)$ of T using a cost and/or effectiveness measure $F : S \rightarrow \mathbb{R}$, and then select the subset of tests that scores best on these measures. Because setting $S := \mathcal{P}(T)$ results in an exponential search space, search-based approaches are often combined with other methods – for example generating the subsets by applying clustering methods using a similarity measure.

$$T_s = \arg \max_{s_i \in S} F(s_i)$$

Methods for test suite reduction based on searching are presented in [Cou+13], and CANDO.

3.3.4 Hybrids and Other Methods

Hybrid methods for test suite reduction combine different approaches to test suite reduction to achieve greater performance or speed up the reduction process. For example, search-based test suite reduction approaches can be enhanced by creating the subsets S of T by applying different clustering algorithms.

Other methods that can not be easily classified into any of those three categories have been proposed in the literature, for example Integer Linear Programming-based solutions, which are used for multi-objective optimization (e.g. optimizing for high coverage, low running times and low memory usage of tests).

4 Approach

4.1 Replicating “Scalable Approaches for Test Suite Reduction”

This paper attempts to replicate the findings in [Cru+19], using different test data and adding a further algorithm as a baseline. [Cru+19] builds on [Mir+18] (which introduced the FAST family), and adds two new algorithms to the FAST family.

We use the code from the original paper, published online at <https://github.com/ICSE19-FAST/FAST-R>, with some slight modifications to fix faults discovered in the original code.

4.2 Implemented Algorithms

[Cru+19] compares 7 different methods of test-suite reduction. 4 of those (**FAST++**, **FAST-all**, **FAST-CS** and **FAST-pw**) are in the FAST family (first introduced in [Mir+18]), which uses clustering techniques to find representative test cases. 2 other algorithms are taken from the similarly clustering-based ART family, first developed in [Che+10]. The last reduction method examined in [Cru+19] is the Greedy Additional (**GA**) algorithm developed in [Rot+01], which is included because "for its simplicity and effectiveness [it] is often considered as a baseline."

This paper also includes a random selection algorithm as a baseline, as recommended by [Kha+18].

4.2.1 FAST

The FAST family is a clustering-based family of algorithms for test suite reduction ([Mir+18], [Cru+19]).

The FAST family is a collection of 4 clustering based TSR algorithms that work both with adequate and inadequate TSR problems, especially for large test suites.

FAST++

The FAST++ algorithm starts with a preparation phase: the tests from T are transformed into points in a vector space by treating each token (e.g. character) of the test case as a dimension, with the value of that dimension being the value of the token at the position (e.g. the value of the n th character), with the components being "weighted according to [a] term-frequency scheme, i.e., the weights are equal to the frequency of the corresponding terms" ([Cru+19]).

Since dealing with high-dimensional vector spaces is computationally costly (e.g. when computing the Euclidean distance), the algorithm performs a random projection into a lower-dimensional vector-space that nonetheless still mostly preserves the pairwise distances of the vectors (in this case a sparse random projection ([Ach03])).

The second phase of FAST++ is executing the k-means++ algorithm [AV06] on the resulting vectors, with k in the inadequate case being the budget of the reduction, and in the adequate case being a variable that is incremented until the requirements are met. The k-means algorithm is a clustering algorithm that finds k clusters of vectors in a high-dimensional space, i.e. minimizing the distance between points within a cluster, by iteratively assigning points to the nearest mean (distance being measured in squared Euclidean distances) and recalculating means until a fixed point is reached. The k-means++ algorithm only differs from k-means in the method for choosing initial centers of clusters: While k-means selects values at random, k-means++ selects the initial by computing the center of all points, and then selecting k other centers by sampling points using a distribution proportional to the distance of the points to the global mean. This both increases speed and gives guarantees that the solution is $\mathcal{O}(\log k)$ competitive to the optimal solution.

After computing the k clusters, FAST++ returns the vectors closest to the centers of the k clusters (i.e. the tests most representative for those clusters).

FAST-CS

The FAST-CS algorithm has the same preparation phase as FAST++: test cases are vectorized, and the resulting vectors are projected into a lower-dimensional vectorspace.

The second phase is different, as it attempts to cluster the set of points by finding a coresot, a set of points that approximate the shape of the set of vectors.

This is achieved by importance sampling: "All points have nonzero probability of being sampled, but points that are far from the center of the dataset (potentially good centers for clustering) are sampled with higher probability."

The metric used for importance sampling is as follows:

$$Q(t) \leftarrow \frac{1}{2|T|} + \frac{d(P(t), \mu)^2}{\sum_{t' \in P} d(P(t'), \mu)^2}$$

where μ is the mean of the whole dataset, T is the test suite, d is a distance metric (in this case the Euclidean metric), and P is the random projection of the test suite.

The points are sampled without replacement until either the coverage is adequate or the budget has been reached.

FAST-pw

FAST-pw was first introduced in [Mir+18] as a test suite prioritization algorithm. It attempts to maximize the Jaccard distance between early test cases (the Jaccard distance between two sets A and B being $JD(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}$).

The algorithm uses three different concepts: *shingling*, *minhashes* (and *minhash signatures*), and *locality-sensitive hashing*.

Shingling is used as a method for transforming a test case into a set, and is therefore only used in the case where coverage information is not available. Given a string of S , a k -shingle is the set of all substrings of S of length k . For example, the 7-shingle of "bananas" is just the set $\{"bananas"\}$, the 5-shingle of "bananas" is $\{"banan", "anana", "nanas"\}$. [Mir+18] states that "if two documents are similar they will have many shingles in common".

However, as [Mir+18] observes, sets of shingles (or of coverage information) can be very large.

Minhashing is a method for deriving compact representations of sets. This is achieved by creating a list of hash functions $H = \{h_1, \dots, h_k\}$, and for each test case T (a set of shingles or of coverage information) calculating the list $[\arg \min_{t \in T} h_1(t), \dots, \arg \min_{t \in T} h_k(t)]$ (which is called the signature of the set). Less formally, the minhash at position i is the minimal value of the hash function for any element of the set.

The Jaccard distance of two sets can be estimated by calculating the number of positions on which the minhash signatures s_1, s_2 of the two sets agree:

$$\text{EstimateJD}(s_1, s_2) = 1 - \frac{|\{i | i \in 1..k, s_1(i) = s_2(i)\}|}{|s_1|}$$

Locality-sensitive hashing is a further technique to make runtimes shorter. Let $S = \{s_1, \dots, s_n\}$ be a set of minhash signatures of all tests. Now a matrix $M \in \mathbb{R}^{n \times k}$ is created, by using the minhash signatures as columns. The rows of this matrix are now divided into bands of length r (in the example 2):

b_n is the n th band, r_n is the n th row, and s_n is the n th signature.

	s_1	s_2	s_3	s_4	s_5	s_6	s_7	s_8	
r_1	1	16	9	5	5	1	2	6	b_1
r_2	2	19	16	18	0	9	5	5	
	8	10	9	14	1	3	11	2	lsh_1
r_3	10	8	4	11	18	9	16	7	b_2
r_4	12	14	3	5	8	3	1	16	
	3	5	1	6	17	13	12	7	lsh_2
r_5	10	11	5	18	14	6	14	7	b_3
r_6	18	5	2	13	7	19	9	11	
	7	6	6	13	6	14	12	16	lsh_3

Figure 4.1: Locality-Sensitive Signature Hash Matrix

The columns within the bands are hashed into lsh values. If two lsh values correspond (e.g. in the case where the columns with values 11 and 5 are both hashed into 6), they are put into the same candidate set.

When the hash of the values in different columns of two bands is the same (the two minhashes are in the same bucket), the two test cases the minhashes belong to are added to the confusingly named candidate set (tests in the candidate set are excluded from selection). The candidate set contains tests that have a Jaccard distance over a below threshold s , where $s \approx 1 - (r/n)^{1/r}$.

The algorithm follows different steps:

1. Calculate the minhashes
2. Calculate the candidate set C_s
3. Select a test from the set of remaining tests $T \setminus C_s$
4. Remove that test from the remaining tests T and add it to the set of selected tests S

5. If the condition is not fulfilled (budget not fulfilled in the inadequate case, coverage not achieved in the adequate case), go to step 2
6. Return S

The function that selects the test case from the remaining test cases maximizes the Jaccard distance of the returned test and the so far selected test cases S (where M is a function that produces minhashes, and Estimate:

$$\arg \max_{c \in T \setminus C_s} \text{EstimateJD}(M(S), M(c))$$

FAST-all

FAST-all is very similar to FAST-pw. The only difference lies in the selection function: While FAST-pw chooses based on Jaccard distance, FAST-all selects all of $T \setminus C_s$, simply choosing all non-candidate tests. This is a special case for the originally described algorithm, which allowed to select a random sample from the non-candidate tests with a given size (*one*, *log*, *sqr*t, *all*).

4.2.2 Random Selection

$\frac{1}{2}$ a page

[Kha+18] gives 12 recommendations on how to execute and evaluate test suite reduction experiments. One of them is to include a baseline algorithm to compare the other algorithms to, such as random selection.

In the inadequate case, random selection simply chooses B random tests from the test set:

```

function RSBUDGET( $TS, B$ )
  return randomsubset( $TS, B$ )
end function

```

In the adequate case, a random element is removed from the original test suite and added to the selected tests until the set of selected tests fulfills all requirements (in this case, covers all coverable lines):

4.2.3 Adaptive Random Testing

The second family of algorithms compared to in this paper is the ART (Adaptive Random Testing) family introduced in [Jia+09] and elaborated on in [Che+10]. These

```
function RSADEQUATE( $TS, cov, B$ )  
   $RS \leftarrow \emptyset$   
   $allcov \leftarrow \bigcup_{t \in TS} cov(t)$   
  while  $\bigcup_{t \in RS} cov(t) \neq allcov$  do  
     $s \leftarrow selectrandom(TS)$   
     $TS \leftarrow TS \setminus \{s\}$   
     $RS \leftarrow RS \cup \{s\}$   
  end while  
  return  $RS$   
end function
```

algorithms were originally developed for test case prioritization (as an alternative to random selection), but can be trivially converted into test suite reduction algorithms by applying the prioritization procedure and then either selecting the first B test cases or so many that they fulfill all requirements.

ART-D

The ART-D algorithm was presented in [Jia+09] (although there only called ART). It starts with a set of tests T and an empty sequence of prioritized tests P . The algorithm works in two stages:

First, tests are selected from T and added to a candidate set C as long as new tests add new coverage information to C (this will not result in the joint coverage of tests in C being the joint coverage of tests in T). This results in the size of candidate sets being variable during execution.

Second, a test c is selected from C and appended to P . The criterion for selecting c is to maximize the distance between all $p \in P$ and c ; [Jia+09] use the Jaccard distance in coverage and examine different methods for comparing the sequence of prioritized test cases to the tests in the candidate set (maximizing the maximum/average/minimum distance between the selected test case and the already prioritized cases (and conclude that maximizing the minimal distance might be the best strategy)).

The selected test case is appended to the sequence of prioritized test cases and removed from T . Unless T is now empty, the first stage is executed again.

ART-F

The ART-F algorithm was introduced in [ZSS12]. It is similar in structure to ART-D, but uses the Manhattan distance instead of the Jaccard distance, and frequency information instead of coverage information, as well as having a fixed-size candidate set. Frequency

information is information about how often a part of the SUT is executed by a test. However, in the presented experiment, since frequency information is not available, coverage information is used as a place-holder, where a test covering a statement is treated as the test executing the statement with frequency 1. The Manhattan distance of two sequences of frequency information is calculated as such:

$$MD(u, v) = \sum_{i=1}^n |u_i - v_i|$$

The stages in ART-F are the same as in ART-D, and it also returns a sequence of prioritized test cases.

4.2.4 Greedy Algorithm

The greedy algorithm used is the greedy additional algorithm from [Rot+01]: In each iteration, it selects the test case with the highest number of yet uncovered program entities.

```

function GANEXT( $RS, TS, cov$ )
  alreadycov  $\leftarrow \bigcup_{t \in RS} cov(t)$ 
  return  $\arg \max_{t \in TS} |cov(t) \setminus alreadycov|$ 
end function

```

In the adequate case, we select greedily until we achieve complete coverage:

```

function GAADEQUATE( $TS, cov$ )
   $RS \leftarrow \emptyset$ 
   $allcov \leftarrow \bigcup_{t \in TS} cov(t)$ 
  while  $\bigcup_{t \in RS} cov(t) \neq allcov$  do:
     $s \leftarrow ganext(RS, TS, cov)$ 
     $TS \leftarrow TS \setminus \{s\}$ 
     $RS \leftarrow RS \cup \{s\}$ 
  end while
  return  $RS$ 
end function

```

In the budget case, we instead select greedily B times.

function GABUDGET(TS, cov, B)

$RS \leftarrow \emptyset$

$allcov \leftarrow \bigcup_{t \in TS} cov(t)$

for doB times

$s \leftarrow ganext(RS, TS, cov)$

$TS \leftarrow TS \setminus \{s\}$

$RS \leftarrow RS \cup \{s\}$

end for

return RS

end function

5 Replication Study

17 pages

5.1 Research Questions

We define five research questions, adapting three from [Cru+19] and adding two further questions.

Research Question 1.1: Can the relative effectiveness in TSR be replicated on other study objects? We care about the size of the reduced test suite, therefore it is valuable by how much test suites are reduced. We use the TSR metric (defined in terms and definitions) to evaluate the magnitude of the reduction. Using the TSR is mostly useful in adequate scenarios, since in inadequate scenarios, the size of the reduced test suite relative to the original test suite is fixed. We rank the algorithms using a Kruskal-Wallis test, taken from [Cru+19]

Research Question 1.2: Can the relative effectiveness in FDL be replicated on other study objects? [Rot+02] finds that test suite reduction can severely compromise fault detection capability. To test whether this is the case, the FDL metric is used to examine the fault detection capabilities for the algorithms given new data. We also rank the methods using a Kruskal-Wallis test.

Research Question 2: Can the relative runtime performance of the different algorithms be replicated? Due to using different datasets (of both larger and smaller size), the absolute runtimes of the different test suite reduction methods can be expected to be different from the original runs. However, the ordering of the methods according to runtime is expected to be the same (that is, if algorithm X was faster than algorithm Y on the old data, we can expect algorithm X to also be faster than algorithm Y on the new data).

Again, we use the Kruskal-Wallis test to rank methods after their runtime performance.

Research Question 3: How much better than random selection are specialized algorithms? [Kha+18] recommends a comparison of the examined technique to a simple baseline. We use their recommendation (random selection) and compare the other approaches to how well they fare against random selection on runtime, TSR and FDL.

(Possibly) Research Question 4: Can the results be reproduced with the original test data?

5.2 Study Design

We approached the research questions by collecting fault data, coverage data and test case source code from projects that had not been used in [Cru+19], and running the code from that paper on the new data. We then generated descriptive statistics (the mean, median, and standard deviation) on the performance (i.e. FDL, TSR, and runtime) of the different algorithms. Furthermore, we ordered the methods by median, and performed a Kruskal-Wallis test to create a statistically significant ordering of the different methods on the metrics evaluated.

5.3 Study Objects

The code bases and test suites selected for this study were small to medium sized open-source Java projects. They all used either JUnit 4 or JUnit 5 as their testing framework, and the tool Maven for building and testing. We used the latest version.

We selected java projects because tools for generating coverage information and mutation-test data using Maven are available and comparatively easy to use. We decided to use open-source projects to make checking our generated data easier, as well as for availability reasons.

The selected projects are both under development and in use, which makes our results relevant to real-life systems.

5.4 Selecting Projects

The projects used, their versions and size are presented in table CANDO (project and test suite size are in lines of code, the version is the 6-digit prefix of the git commit used).

Table 5.1

Project name	Version	Project size	Test-suite size	Number of tests
assertj-core	cb2829	135k	241k	3578
commons-collections	242918	59k	45k	238
commons-lang	6b3f25	51k	40k	181
commons-math	649b13	148k	98k	438
jopt-simple	5a1d72	1.7k	2.7k	145
jsoup	89580c	18k	12k	52

5.5 Study Setup

Testing the performance of the algorithms required three different kinds of information: the contents of the test suites, coverage information and fault information.

5.5.1 Combining Tests Suites

For every examined project, the test suite was converted into a format suitable for the code from [Cru+19] by replacing the newlines from each test with spaces and concatenating the tests into one file (the **black-box file**), such that every line contained one test case.

For some of the projects, test cases were excluded since they failed during the default test run:

Table 5.2

Project name	Test classes excluded
assertj-core	BDDSoftAssertionsTest, SoftAssertionsTest, SoftAssertions_overriding_afterAssertionErrorCollected_Test, SoftAssertionsErrorsCollectedTest, AssertionsMultipleProjectsTest, SoftAssertions_setAfterAssertionErrorCollected_Test, AssertJ-MultipleFailuresErrorMessage_Test
commons-collections	BulkTest
commons-lang	FieldUtilsTest
commons-math	FastMathTest, EvaluationTestValidation
jopt-simple	/
jsoup	/

5.5.2 Generating Coverage Information

Coverage information was generated using jacoco with the testwise mode of the teamscale jacoco agent.

Since the teamscale jacoco agent only supports line coverage, other types of coverage had to be eschewed.

The json files created by the teamscale jacoco agent were converted from JSON into the format used by [Cru+19]: a file containing a list of numbers in each line, with the numbers n_1, \dots, n_i in line n corresponding to the source lines of code in the tested project covered by the test case at line n in the **black-box file**.

5.5.3 Collecting Fault Coverage Information

Fault detection information was not available for the projects used, and was therefore generated by mutation testing using the mutation testing framework pitest (first described in [Bud80]).

The generated mutation test data was converted from XML to the format used in the code of [Cru+19]: a text file containing a list of number per line, the numbers at line n corresponding to different classes for which the test case at line n in the **black-box file** did find faults.

5.6 Results

To obtain the results, we used the code from [Cru+19], and applied the test suite reduction algorithms to new test data (i.e. black-box test suite files and coverage information of 6 open source projects) after minimal modification. These modifications were as follows:

We added a further algorithm for the random selection of test cases.

We fixed three bugs we found in the code: an off-by-one error in loading coverage information, a bug where in small test suites sometimes the candidate set would be empty, and a mistake where a measured variable was not reported correctly for FAST-all. The changes made to the source code can be found at CANDO.

In the budget scenario, we considered budgets between 1% and 30%, with a step increase of 1%. We considered only line coverage (functionally equivalent to statement coverage), and performed 50 measurements of FDL, TSR, preparation and reduction time.

In the adequate scenario, we also considered only line coverage, and made 50 measurements of preparation time, coverage preparation time, and reduction time.

We did not replicate the large-scale scenario.

All measurements were performed under a Ubuntu 20.04 64-bit system, using 6 AMD® Ryzen 5 4500u with radeon graphics processors, and 7.2 GiB of RAM available.

5.6.1 Research Questions

Research Question 1.1 (Can TSR findings be replicated?) Table 5.3 shows the descriptive statistical results (median, standard deviation and ordering for the Kruskal-Wallis test) in the adequate case.

For Java programs with statement coverage in the adequate case, [Cru+19] order the different methods of test suite reduction according to TSR. Their results mostly agree with ours: They rank GA highest, then FAST++ and FAST-pw, then FAST-CS, FAST-all below that, and the ART family lowest; we also rank GA highest, followed by FAST++, then FAST-CS and FAST-pw at the same level, then FAST-all, and the ART family last (on par with random selection). However, the median test suite reduction they report for these methods is much lower than what we observe (at least for the java programs): even GA only achieves a TSR of 12.30, while in our experiment GA achieves a TSR of over 40. However, their results with C programs are more stark than ours, achieving a TSR of more than 90 for GA or FAST++.

Our data shows higher variance in TSR performance on all algorithms.


Table 5.3: Different variables relating to TSR (mdn is the median, σ being the standard deviation, and δ being the ranking in the Kruskal-Wallis test)

Approach	μ	mdn	σ	δ
ART-D	18.421	17.351	8.997	(e)
ART-F	0.945	0.228	1.654	(f)
FAST++	36.116	33.811	13.581	(b)
FAST-all	21.414	18.264	8.536	(d)
FAST-CS	33.683	31.893	12.064	(c)
FAST-pw	33.258	32.718	11.458	(c)
GA	40.599	39.124	14.922	(a)
RS	0.947	0.420	1.491	(f)

Research Question 1.2 Table 5.4 shows the descriptive statistics (mean, median, standard deviation, and ordering according to the Kruskal-Wallis test) for the FDL of different algorithms in both the budget and adequate case.

We compare our findings to FDL performance on statement-covered C program data, since the fault information for the java programs used in [Cru+19] is very sparse (1-3

Figure 5.1: Boxplots for TSR for different algorithms in the adequate case



faults per program), while fault information is much richer for the C programs.

In the budget scenario, we could not detect strong differences in performance in FDL. All algorithms score equally, with the exception of GA dominating every other algorithm and RS being dominated by every other algorithm. Here, again, [Cru+19] come to very similar conclusions: They find that GA is statistically significantly better than all other algorithms, which are in turn approximately equally effective at detecting faults.

Here, we also find that the standard deviations in FDL for our data is higher than the ones in [Cru+19]. Similar to their findings, we see no strong difference in variance for different algorithms.

In the adequate case, our data ranks the algorithms into just two categories: The better-performing category contains algorithms from the ART family, FAST++ and FAST-CS, while the rest of the approaches is sorted into the bucket of worse-performing approaches. [Cru+19] agrees that the ART family is best-performing, but find that FAST-all is second-best, GA being third-best, and FAST-CS, FAST-pw and FAST++ coming last. At the moment, we don't have a hypothesis about why this might be the case. It seems relevant, however, that the median fault detection loss for data was 0 for all algorithms (the mean fault detection loss being non-zero, however), just as the ART

family and FAST-all in [Cru+19]. Perhaps this is caused by single tests both having large coverage and detecting a majority of faults.

This hypothesis is further supported by the observation that the variance in performance on all algorithms is much lower than in the original paper, for all algorithms.

Table 5.4: Different variables relating to FDL (μ being the mean, mdn the median, σ the standard deviation, and δ the ranking in the Kruskal-Wallis test)

Budget					Adequate				
Approach	μ	mdn	σ	δ	Approach	μ	mdn	σ	δ
ART-D	76.414	94.75	34.502	(b)	ART-D	0.129	0.0	0.702	(a)
ART-F	77.72	95.04	35.684	(b)	ART-F	0.057	0.0	0.679	(a)
FAST++	76.49	94.05	35.662	(b)	FAST++	1.122	0.0	2.172	(a)
FAST-all	74.78	92.19	35.269	(b)	FAST-all	0.155	0.0	0.331	(b)
FAST-CS	76.19	93.39	35.493	(b)	FAST-CS	1.144	0.0	2.367	(a)
FAST-pw	71.63	93.82	35.665	(b)	FAST-pw	0.792	0.0	1.131	(b)
GA	62.53	70.45	33.367	(a)	GA	1.432	0.0	2.444	(b)
RS	93.43	97.46	9.305	(c)	RS	0.058	0.0	0.679	(b)

We include the mean in the descriptive statistics to illustrate that while the median adequate reduction has no fault detection loss, it can nonetheless still occur.

Research Question 2

Research Question 3

(Possibly) Research Question 4


5.7 Discussion

Comparison to "Scalable Approaches to Test Suite Reduction"

5.8 Threats to Validity

Although we attempted to prevent it, the results we obtained might suffer from threats to validity.

Figure 5.2: Boxplots for FDL for different algorithms



5.8.1 Construct Validity

Construct validity considers whether the methods are adequate to answer the research questions posed. As in the original paper, we used FDL and TSR as standard metrics, and reported the mean, median and standard deviation of the data as standard statistical properties of the generated data. We used faults generated by mutation testing software instead of real faults, which might be dissimilar from real faults.

5.8.2 Internal Validity

The results we obtained might not be due to actual differences in the performance of the different methods, but due to other factors, such as the usage of the computer used for gathering data for other purposes during that time. To mitigate these effects, we performed every measurement 50 times, and abstained from using the computer during the time of gathering data.

5.8.3 Conclusion Validity

Threats to conclusion validity are present when false positives are found due to considering many different test suite reduction methods and variables. This paper performs 4 statistical tests, thus we judge a significance level of 5% to be sufficient.

5.8.4 External Validity

The projects we selected are all small to medium-sized projects, and therefore the results we obtained might not generalize to large-scale software systems. While one of the projects we obtained test data from (assertj-core) was larger than the projects examined in [Cru+19], it was not analyzed separately from the other data. However, since the projects we used were of similar size as the ones in [Cru+19], and our interest was to replicate their findings, we regard this consideration as less relevant.

6 Future Work

1 page

7 Summary

2 pages

List of Figures

4.1	Locality-Sensitive Signature Hash Matrix	11
5.1	TSR boxplots, adequate	21
5.2	FDL boxplots	23

List of Tables

5.3	TSR statistical results, adequate	20
5.4	FDL statistical results	22

Bibliography

- [Ach03] D. Achlioptas. "Database-friendly random projections: Johnson-Lindenstrauss with binary coins." In: *Journal of computer and System Sciences* 66.4 (2003), pp. 671–687.
- [AV06] D. Arthur and S. Vassilvitskii. *k-means++: The advantages of careful seeding*. Tech. rep. Stanford, 2006.
- [Bud80] T. A. Budd. "Mutation analysis of program test data[Ph. D. Thesis]." In: (1980).
- [Che+10] T. Y. Chen, F.-C. Kuo, R. G. Merkel, and T. Tse. "Adaptive random testing: The art of test case diversity." In: *Journal of Systems and Software* 83.1 (2010), pp. 60–66.
- [Cou+13] A. Coutinho, E. G. Cartaxo, P. D. Machado, and C. G. SPLab-UFCG. "Test suite reduction based on similarity of test cases." In: *7st Brazilian workshop on systematic and automated software testing—CBSOft*. Vol. 2013. 2013.
- [Cru+19] E. Cruciani, B. Miranda, R. Verdecchia, and A. Bertolino. "Scalable approaches for test suite reduction." In: *2019 IEEE/ACM 41st International Conference on Software Engineering (ICSE)*. IEEE. 2019, pp. 419–429.
- [GJ79] M. R. Garey and D. S. Johnson. *Computers and intractability*. Vol. 174. freeman San Francisco, 1979.
- [Jia+09] B. Jiang, Z. Zhang, W. K. Chan, and T. Tse. "Adaptive random test case prioritization." In: *2009 IEEE/ACM International Conference on Automated Software Engineering*. IEEE. 2009, pp. 233–244.
- [Kha+16] S. U. R. Khan, S. P. Lee, R. W. Ahmad, A. Akhunzada, and V. Chang. "A survey on Test Suite Reduction frameworks and tools." In: *International Journal of Information Management* 36.6 (2016), pp. 963–975.
- [Kha+18] S. U. R. Khan, S. P. Lee, N. Javaid, and W. Abdul. "A systematic review on test suite reduction: Approaches, experiment's quality evaluation, and guidelines." In: *IEEE Access* 6 (2018), pp. 11816–11841.

- [Mir+18] B. Miranda, E. Cruciani, R. Verdecchia, and A. Bertolino. "Fast approaches to scalable similarity-based test case prioritization." In: *2018 IEEE/ACM 40th International Conference on Software Engineering (ICSE)*. IEEE. 2018, pp. 222–232.
- [Rot+01] G. Rothermel, R. H. Untch, C. Chu, and M. J. Harrold. "Prioritizing test cases for regression testing." In: *IEEE Transactions on software engineering* 27.10 (2001), pp. 929–948.
- [Rot+02] G. Rothermel, M. J. Harrold, J. Von Ronne, and C. Hong. "Empirical studies of test-suite reduction." In: *Software Testing, Verification and Reliability* 12.4 (2002), pp. 219–249.
- [YH12] S. Yoo and M. Harman. "Regression testing minimization, selection and prioritization: a survey." In: *Software testing, verification and reliability* 22.2 (2012), pp. 67–120.
- [ZSS12] Z. Q. Zhou, A. Sinaga, and W. Susilo. "On the fault-detection capabilities of adaptive random test case prioritization: Case studies with large test suites." In: *2012 45th Hawaii International Conference on System Sciences*. IEEE. 2012, pp. 5584–5593.